

WES 237A: Introduction to Embedded System Design (Winter 2026)

Lab 2: Process and Thread

Due: 1/19/2026 11:59pm

In order to report and reflect on your WES 237A labs, please complete this Post-Lab report by the end of the weekend by submitting the following 2 parts:

- Upload your lab 2 report, composed by a single PDF that includes your in-lab answers to the bolded questions in the Google Doc Lab and your Jupyter Notebook code.
- Answer two short essay-like questions on your Lab experience.

All responses should be submitted to Canvas. Please also be sure to push your code to your git repo as well.

Create Lab2 Folder

1. Create a new folder on your PYNQ jupyter home and rename it 'Lab2'

Shared C++ Library

1. In 'Lab2', create a new text file (New -> Text File) and rename it to 'main.c'
2. Add the following code to 'main.c':

```
#include <unistd.h>
```

```
int myAdd(int a, int b){  
    sleep(1);  
    return a+b;  
}
```

3. Following the function above, write another function to multiply two integers together. Copy your code below.

```
int myMulti(int a, int b){  
    sleep(1);  
    return a*b;  
}
```

4. Save main.c
5. In Jupyter, open a terminal window (New -> Terminal) and *change directories* (cd) to 'Lab2' directory.

```
$ cd Lab2
```

6. Compile your 'main.c' code as a shared library.

```
$ gcc -c -Wall -Werror -fpic main.c  
$ gcc -shared -o libMyLib.so main.o
```

7. Download 'ctypes_example.ipynb' from [here](#) and upload it to the Lab2 directory.
8. Go through each of the code cells to understand how we interface between Python and our C code

- 9. Write another Python function to wrap your multiplication function written above in step 3. Copy your code below.**

```
def multiC(a,b):  
    return _libInC.myMulti(a,b)
```

To summarize, we created a C shared library and then called the C function from Python

Multiprocessing

1. Download ‘multiprocess_example.ipynb’ from [here](#) and upload it to your ‘Lab2’ directory.
2. Go through the documentation (and comments) and answer the following question
 - a. **Why does the ‘Process-#’ keep incrementing as you run the code cell over and over?**

New processes are just being added to the RAM in the working environment.

- b. **Which line assigns the processes to run on a specific CPU?**

```
os.system("taskset -p -c {} {}".format(0, p1.pid))
```

3. In ‘main.c’, change the ‘sleep()’ command and recompile the library with the commands above. Also, reload the Jupyter notebook with the ⌂ symbol and re-run all cells. Play around with different sleep times for both functions.
 - a. **Explain the difference between the results of the ‘Add’ and ‘Multiply’ functions and when the processes are finished.**

The main difference between the two functions is that one adds and the other multiplies. Multiply always finishes a slight bit later, but that's because it is called after Add.

4. Continue to the lab work section. Here we are going to do the following
 - a. Create a multiprocessing array object with 2 entries of integer type.
 - b. Launch 1 process to compute addition and 1 process to compute multiplication.
 - c. Assign the results to separate positions in the array.
 - i. Process 1 (add) is stored in index 0 of the array (array[0])
 - ii. Process 2 (mult) is stored in index 1 of the array (array[1])
 - d. Print the results from the array.
 - e. **There are 4 TODO comments that must be completed**
5. Answer the following question
 - a. **Explain, in your own words, what shared memory is in relation to the code in this exercise.**

The shared memory in this instance is the use of the same array to process different arithmetic operations on different cores. ‘Add’ is on returnValues[0] and ‘Multiply’ is happening on returnValues[1], which are adjacent to each other in memory.

Threading

1. Download ‘threading_example.ipynb’ from [here](#) and upload it into your ‘Lab2’ directory.
2. Go through the documentation and code for ‘Two threads, single resource’ and answer the following questions
 - a. **What line launches a thread and what function is the thread executing?**

t.start(), executes worker_t

- b. **What line defines a mutual resource? How is it accessed by the thread function?**

fork = threading.Lock(), it is accessed by passing in “fork” in args

3. Answer the following question about the ‘Two threads, two resources’ section.
 - a. **Explain how this code enters a deadlock.**

Thread0 has resource1 and Thread1 has resource0. The threads are waiting for the other to release their respective resources to proceed. However, since neither of them end up doing that, they both get stuck and we end up with a deadlock.

4. Complete the code using the non-blocking acquire function.
 - a. **What is the difference between ‘blocking’ and ‘non-blocking’ functions?**

“Blocking” means code execution will halt until a condition is met. “Non-blocking” means the code will keep going regardless if that condition is met.

5. BONUS:

Can you explain why this is used in the ‘Two threads, two resources’ section:

```
if using_resource0:  
    _l0.release()  
if using_resource1:  
    _l1.release()
```

```
1 #include <unistd.h>
2
3 int myAdd(int a, int b){
4     sleep(10);
5     return a+b;
6 }
7
8 int myMulti(int a, int b){
9     sleep(10);
10    return a*b;
11 }
```


ctypes

The following imports ctypes interface for Python

```
In [ ]: import ctypes
```

Now we can import our shared library

```
In [ ]: _libInC = ctypes.CDLL('./libMyLib.so')
```

Let's call our C function, myAdd(a, b).

```
In [ ]: print(_libInC.myAdd(3, 5))
```

This is cumbersome to write, so let's wrap this C function in a Python function for ease of use.

```
In [ ]: def addC(a,b):
         return _libInC.myAdd(a,b)
```

Usage example:

```
In [ ]: print(addC(10, 202))
```

Multiply

Following the code for your add function, write a Python wrapper function to call your C multiply code

```
In [ ]: def multiC(a,b):
         return _libInC.myMulti(a,b)
```

multiprocessing

importing required libraries and our shared library

```
In [1]: import ctypes
import multiprocessing
import os
import time
```

```
In [2]: _libInC = ctypes.CDLL('./libMyLib.so')
```

Here, we slightly adjust our Python wrapper to calculate the results and print it. There is also some additional casting to ensure that the result of the *libInC.myAdd()* is an int32 type.

```
In [3]: def addC_print(_i, a, b, time_started):
    val = ctypes.c_int32(_libInC.myAdd(a, b)).value #cast the result to int32
    end_time = time.time()
    print('CPU_{0} Add: {1} in {2}'.format(_i, val, end_time - time_started))

def multiC_print(_i, a, b, time_started):
    val = ctypes.c_int32(_libInC.myMulti(a, b)).value #cast the result to int32
    end_time = time.time()
    print('CPU_{0} Multiply: {1} in {2}'.format(_i, val, end_time - time_started))
```

Now for the fun stuff.

The multiprocessing library allows us to run simultaneous code by utilizing multiple processes. These processes are handled in separate memory spaces and are not restricted to the Global Interpreter Lock (GIL).

Here we define two processes, one to run the *addC_print* and another to run the *multiC_print()* wrappers.

Next we assign each process to be run on different CPUs

```
In [4]: procs = [] # a future list of all our processes

# Launch process1 on CPU0
p1_start = time.time()
p1 = multiprocessing.Process(target=addC_print, args=(0, 3, 5, p1_start))
os.system("taskset -p -c {} {}".format(0, p1.pid)) # taskset is an os command
p1.start() # start the process
procs.append(p1)

# Launch process2 on CPU1
p2_start = time.time()
p2 = multiprocessing.Process(target=multiC_print, args=(1, 3, 5, p2_start))
os.system("taskset -p -c {} {}".format(1, p2.pid)) # taskset is an os command
p2.start() # start the process
procs.append(p2)

p1Name = p1.name # get process1 name
p2Name = p2.name # get process2 name

# Here we wait for process1 to finish then wait for process2 to finish
p1.join() # wait for process1 to finish
print('Process 1 with name, {}, is finished'.format(p1Name))

p2.join() # wait for process2 to finish
print('Process 2 with name, {}, is finished'.format(p2Name))

taskset: invalid PID argument: 'None'
taskset: invalid PID argument: 'None'

CPU_0 Add: 8 in 10.07143235206604
CPU_1 Multiply: 15 in 10.13064694404602
Process 1 with name, Process-1, is finished
Process 2 with name, Process-2, is finished
```

Return to 'main.c' and change the amount of sleep time (in seconds) of each function.

For different values of sleep(), explain the difference between the results of the 'Add' and 'Multiply' functions and when the Processes are finished.

Lab work

One way around the GIL in order to share memory objects is to use multiprocessing objects. Here, we're going to do the following.

1. Create a multiprocessing array object with 2 entries of integer type.
2. Launch 1 process to compute addition and 1 process to compute multiplication.
3. Assign the results to separate positions in the array.
 - A. Process 1 (add) is stored in index 0 of the array (array[0])
 - B. Process 2 (mult) is stored in index 1 of the array (array[1])
4. Print the results from the array.

Thus, the multiprocessing Array object exists in a *shared memory* space so both processes

can access it.

Array documentation:

<https://docs.python.org/2/library/multiprocessing.html#multiprocessing.Array> (<https://docs.python.org/2/library/multiprocessing.html#multiprocessing.Array>)

typecodes/types for Array:

'c': ctypes.c_char

'b': ctypes.c_byte

'B': ctypes.c_ubyte

'h': ctypes.c_short

'H': ctypes.c_ushort

'i': ctypes.c_int

'I': ctypes.c_uint

'l': ctypes.c_long

'L': ctypes.c_ulong

'f': ctypes.c_float

'd': ctypes.c_double

Try to find an example

You can use online resources to find an example for how to use multiprocessing Array

```
In [5]: def addC_no_print(_i, a, b, returnValues):
    """
    Params:
        _i : Index of the process being run (0 or 1)
        a, b : Integers to add
        returnValues : Multiprocessing array in which we will store the results
    """
    val = ctypes.c_int32(_libInC.myAdd(a, b)).value
    # TODO: add code here to pass val to correct position returnValues
    returnValues[_i] = val

def multiC_no_print(_i, a, b, returnValues):
    """
    Params:
        _i : Index of the process being run (0 or 1)
        a, b : Integers to multiply
        returnValues : Multiprocessing array in which we will store the results
    """
    val = ctypes.c_int32(_libInC.myMulti(a, b)).value
    # TODO: add code here to pass val to correct position of returnValues
    returnValues[_i] = val

procs = []

# TODO: define returnValues here. Check the multiprocessing docs to see
# about initializing an array object for 2 processes.
# Note the data type that will be stored in the array

returnValues = multiprocessing.Array('i', 2)

p1 = multiprocessing.Process(target=addC_no_print, args=(0, 3, 5, returnValues))
os.system("taskset -p -c {} {}".format(0, p1.pid)) # taskset is an os command
p1.start() # start the process
procs.append(p1)

p2 = multiprocessing.Process(target=multiC_no_print, args=(1, 3, 5, returnValues))
os.system("taskset -p -c {} {}".format(1, p2.pid)) # taskset is an os command
p2.start() # start the process
procs.append(p2)

# Wait for the processes to finish
for p in procs:
    pName = p.name # get process name
    p.join() # wait for the process to finish
    print('{} is finished'.format(pName))

# TODO print the results that have been stored in returnValues
#for i in returnValues:
#    print(returnValues[:])
```

```
taskset: invalid PID argument: 'None'
taskset: invalid PID argument: 'None'
```

```
Process-3 is finished  
Process-4 is finished  
[8, 15]
```

In []:

threading

importing required libraries and programing our board

```
In [1]: import threading  
import time  
from pynq.overlays.base import BaseOverlay  
base = BaseOverlay("base.bit")
```

Two threads, single resource

Here we will define two threads, each responsible for blinking a different LED light. Additionally, we define a single resource to be shared between them.

When thread0 has the resource, led0 will blink for a specified amount of time. Here, the total time is 50×0.02 seconds = 1 second. After 1 second, thread0 will release the resource and will proceed to wait for the resource to become available again.

The same scenario happens with thread1 and led1.

```
In [2]: def blink(t, d, n):
    """
        Function to blink the LEDs
    Params:
        t: number of times to blink the LED
        d: duration (in seconds) for the LED to be on/off
        n: index of the LED (0 to 3)
    """
    for i in range(t):
        base.leds[n].toggle()
        time.sleep(d)
    base.leds[n].off()

def worker_t(_l, num):
    """
        Worker function to try and acquire resource and blink the LED
    _l: threading lock (resource)
    num: index representing the LED and thread number.
    """
    for i in range(4):
        using_resource = _l.acquire(True)
        print("Worker {} has the lock".format(num))
        blink(50, 0.02, num)
        _l.release()
        time.sleep(0) # yeild
    print("Worker {} is done.".format(num))

# Initialize and launch the threads
threads = []
fork = threading.Lock()
for i in range(2):
    t = threading.Thread(target=worker_t, args=(fork, i))
    threads.append(t)
    t.start()

for t in threads:
    name = t.getName()
    t.join()
    print('{} joined'.format(name))
```

Worker 0 has the lock

```
/tmp/ipykernel_12055/1470821917.py:37: DeprecationWarning: getName()
is deprecated, get the name attribute instead
    name = t.getName()
```

```
Worker 0 has the lock  
Worker 0 has the lock  
Worker 0 has the lock  
Worker 0 is done. Worker 1 has the lock
```

```
Thread-5 (worker_t) joined  
Worker 1 has the lock  
Worker 1 has the lock  
Worker 1 has the lock  
Worker 1 is done.  
Thread-6 (worker_t) joined
```

Two threads, two resource

Here we examine what happens with two threads and two resources trying to be shared between them.

The order of operations is as follows.

The thread attempts to acquire resource0. If it's successful, it blinks 50 times x 0.02 seconds = 1 second, then attempts to get resource1. If the thread is successful in acquiring resource1, it releases resource0 and proceeds to blink 5 times for 0.1 second = 0.5 second.

```
In [12]: def worker_t(_l0, _l1, num):
    """
        Worker function to try and acquire resource and blink the LED
        _l0: threading lock0 (resource0)
        _l1: threading lock1 (resource1)
        num: index representing the LED and thread number.
        init: which resource this thread starts with (0 or 1)
    """
    using_resource0 = False
    using_resource1 = False

    for i in range(4):
        using_resource0 = _l0.acquire(True)
        if using_resource1:
            _l1.release()
        print("Worker {} has lock0".format(num))
        blink(50, 0.02, num)

        using_resource1 = _l1.acquire(True)
        if using_resource0:
            _l0.release()
        print("Worker {} has lock1".format(num))
        blink(5, 0.1, num)

        time.sleep(0) # yeild

        if using_resource0:
            _l0.release()
        if using_resource1:
            _l1.release()

    print("Worker {} is done.".format(num))

# Initialize and launch the threads
threads = []
fork = threading.Lock()
fork1 = threading.Lock()
for i in range(2):
    t = threading.Thread(target=worker_t, args=(fork, fork1, i))
    threads.append(t)
    t.start()

for t in threads:
    name = t.getName()
    t.join()
    print('{} joined'.format(name))
```

Worker 0 has lock0

/tmp/ipykernel_12055/855109071.py:44: DeprecationWarning: getName() is deprecated, get the name attribute instead
name = t.getName()

```
Worker 0 has lock1Worker 1 has lock0

Worker 1 has lock1
Worker 0 has lock0
Worker 0 has lock1Worker 1 has lock0

Worker 1 has lock1Worker 0 has lock0

Worker 0 has lock1Worker 1 has lock0

Worker 1 has lock1Worker 0 has lock0

Worker 0 has lock1Worker 1 has lock0

Worker 0 is done.
Thread-25 (worker_t) joined

Exception in thread Thread-26 (worker_t):
Traceback (most recent call last):
  File "/usr/lib/python3.10/threading.py", line 1009, in _bootstrap_inner
    self.run()
  File "/usr/lib/python3.10/threading.py", line 946, in run
    self._target(*self._args, **self._kwargs)
  File "/tmp/ipykernel_12055/855109071.py", line 20, in worker_t
    RuntimeError: release unlocked lock

Thread-26 (worker_t) joined
```

You may have noticed (even before running the code) that there's a problem! What happens when thread0 has resource1 and thread1 has resource0! Each is waiting for the other to release their resource in order to continue.

This is a **deadlock**. Adjust the code to prevent a deadlock. Write your code below:

```
In [20]: # TODO: Write your code here
def worker_t(_l0, _l1, num):
    """
    Worker function to try and acquire resource and blink the LED
    _l0: threading lock0 (resource0)
    _l1: threading lock1 (resource1)
    num: index representing the LED and thread number.
    init: which resource this thread starts with (0 or 1)
    """
    using_resource0 = False
    using_resource1 = False

    for i in range(4):
        if using_resource1:
            _l1.release()
            using_resource1 = False
        using_resource0 = _l0.acquire(True)
        print("Worker {} has lock0".format(num))
        blink(50, 0.02, num)

        if using_resource0:
            _l0.release()
            using_resource0 = False
        using_resource1 = _l1.acquire(True)
        print("Worker {} has lock1".format(num))
        blink(5, 0.1, num)

        time.sleep(0) # yeild

    if using_resource0:
        _l0.release()
    if using_resource1:
        _l1.release()

    print("Worker {} is done.".format(num))

# Initialize and launch the threads
threads = []
fork = threading.Lock()
fork1 = threading.Lock()
for i in range(2):
    t = threading.Thread(target=worker_t, args=(fork, fork1, i))
    threads.append(t)
    t.start()

for t in threads:
    name = t.getName()
    t.join()
    print('{} joined'.format(name))
```

Worker 0 has lock0

/tmp/ipykernel_12055/3214970483.py:47: DeprecationWarning: getName() is deprecated, get the name attribute instead
name = t.getName()

```
Worker 0 has lock1Worker 1 has lock0  
  
Worker 1 has lock1  
Worker 0 has lock0  
Worker 0 has lock1Worker 1 has lock0  
  
Worker 1 has lock1Worker 0 has lock0  
  
Worker 0 has lock1Worker 1 has lock0  
  
Worker 1 has lock1Worker 0 has lock0  
  
Worker 0 has lock1Worker 1 has lock0  
  
Worker 0 is done.  
Thread-39 (worker_t) joined  
Worker 1 has lock1  
Worker 1 is done.  
Thread-40 (worker_t) joined
```

Also, write an explanation for what you did above to solve the deadlock problem.

Your answer:

Bonus: Can you explain why this is used in the worker_t routine?

```
if using_resource0:  
    _l0.release()  
if using_resource1:  
    _l1.release()
```

Hint: Try commenting it out and running the cell, what do you observe?

Non-blocking Acquire

In the above code, when *.acquire(True)* was used, the thread stopped executing code and waited for the resource to be acquired. This is called **blocking**: stopping the execution of code and waiting for something to happen. Another example of **blocking** is if you use *input()* in Python. This will stop the code and wait for user input.

What if we don't want to stop the code execution? We can use non-blocking version of the *acquire()* function. In the code below, *resource_available* will be True if the thread currently has the resource and False if it does not.

Complete the code to and print and toggle LED when lock is not available.

```
In [19]: def blink(t, d, n):
    for i in range(t):
        base.leds[n].toggle()
        time.sleep(d)

    base.leds[n].off()

def worker_t(_l, num):
    for i in range(10):
        resource_available = _l.acquire(False) # this is non-blocking .
        if resource_available:

            # write code to:
            # print message for having the key
            # blink for a while
            # release the key
            # give enough time to the other thread to grab the key
            print("Worker {} took lock!".format(num))
            blink(50, 0.02, num)
            _l.release()
            print("Lock released.")
            time.sleep(1)

    else:
        # write code to:
        # print message for waiting for the key
        # blink for a while with a different rate
        # the timing between having the key + yield and waiting for
        print("Hold on! Waiting for lock!.")
        blink(5, 0.1, num)

    print('Worker {} is done.'.format(num))

threads = []
fork = threading.Lock()
for i in range(2):
    t = threading.Thread(target=worker_t, args=(fork, i))
    threads.append(t)
    t.start()

for t in threads:
    name = t.getName()
    t.join()
    print('{} joined'.format(name))
```

Worker 0 took lock!
Hold on! Waiting for lock!.

/tmp/ipykernel_12055/1181751919.py:42: DeprecationWarning: getName() is deprecated, get the name attribute instead
name = t.getName()

```
Hold on! Waiting for lock!.
Hold on! Waiting for lock!.
Lock released.
Worker 1 took lock!
Hold on! Waiting for lock!.
Hold on! Waiting for lock!.
Lock released.
Worker 0 took lock!
Hold on! Waiting for lock!.
Hold on! Waiting for lock!.
Lock released.
Worker 1 took lock!
Hold on! Waiting for lock!.
Hold on! Waiting for lock!.
Lock released.
Worker 0 took lock!
Hold on! Waiting for lock!.
Hold on! Waiting for lock!.
Lock released.
Worker 1 took lock!
Hold on! Waiting for lock!.
Hold on! Waiting for lock!.
Lock released.
Worker 0 took lock!
Worker 1 is done.
Lock released.
Worker 0 is done.
Thread-37 (worker_t) joined
Thread-38 (worker_t) joined
```

In []: